# CS 589: Text Mining and Information Retrieval

# **Lecture 3 - Complete Study Guide**

# IR Evaluation & (Pseudo)-Relevance Feedback

Course: CS 589 - Text Mining and Information Retrieval

**Institution:** Stevens Institute of Technology

Lecture: 3 - IR Evaluation & Relevance Feedback

Topics: Precision/Recall, MAP, MRR, NDCG, TREC, Pooling, A/B Testing, Rocchio Feedback

Prepared for: Midterm Examination

# Study Guide Overview

#### What's in This Guide:

This comprehensive study guide covers all material from Lecture 3 of CS 589, focusing on Information Retrieval evaluation metrics and relevance feedback techniques. The guide includes:

# Complete Topic Coverage

- All evaluation metrics with worked examples
- Evaluation methodologies (TREC, Pooling, Online)
- Relevance feedback techniques (Rocchio, Pseudo-relevance)

#### Exam-Focused Content

- Step-by-step calculation examples
- Common pitfalls and how to avoid them
- Practice problems with detailed solutions
- Formula quick reference sheet

# Learning Resources

- Video tutorials (with time stamps)
- Python code examples
- Interactive demos
- Additional reading materials

# 🔽 Exam Preparation

• Pre-exam checklist

- Time management strategies
- Last-minute review guide

#### **How to Use This Guide:**

1. First Pass: Read through sections 1-4 to understand all concepts

2. A Practice: Complete the practice problems (section 5) by hand

3. **Reinforce:** Watch recommended videos for difficult topics

4. Review: Use the quick reference sheet and checklist before exam

5. **© Final Review:** Follow the "30 Minutes Before Exam" guide

**Estimated Study Time:** 4-6 hours for complete mastery



# **Table of Contents**

# **Quick Navigation**

Section	Торіс	Page Reference
1	Quick Review: Lecture 2 Concepts	RSJ, BM25, Language Models
2	Evaluation Metrics	Precision, Recall, MAP, MRR, NDCG
3	Evaluation Methodologies	TREC, Pooling, A/B Testing
4	Relevance Feedback	Rocchio, Pseudo-relevance
5	Practice Problems	Worked Examples
6	Additional Resources	Videos, Code, Links
4	•	•

#### **Detailed Contents**

#### 1. Quick Review: Lecture 2 Concepts

- RSJ Model (Robertson & Spark Jones)
- BM25 Model
- Language Model-Based Retrieval

#### 2. Evaluation Metrics ( Most Important)

- 2.1 Precision and Recall
- 2.2 Mean Average Precision (MAP)
- 2.3 Mean Reciprocal Rank (MRR)
- 2.4 Normalized Discounted Cumulative Gain (NDCG)

# 3. Evaluation Methodologies

• 3.1 The Cranfield Experiment (1958)

- 3.2 Pooling Strategy
- 3.3 Text REtrieval Conference (TREC)
- 3.4 Online Evaluation & A/B Testing
  - Position Bias
  - User Click Logs
  - Decoy Effects
  - A/B Testing & Statistical Tests
  - Multi-Armed Bandit
  - Interleaving

# 4. Relevance Feedback

- 4.1 Motivation
- 4.2 Rocchio Feedback (Vector Space Model)
- 4.3 Pseudo-Relevance Feedback
- 4.4 Feedback Language Model
- 4.5 Query Expansion/Reformulation

# 5. Practice Problems

- Problem 1: Calculate MAP
- Problem 2: Calculate NDCG@5
- Problem 3: Calculate MRR
- Problem 4: Feedback Language Model
- Problem 5: Understanding IDCG

# 6. Additional Resources

- Recommended Reading
- Video Tutorials
- Practice Tools & Code
- Comparison Tables
- Quick Reference Sheet
- Pre-Exam Checklist

# **Symbol Legend**

Throughout this guide, you'll see these symbols to help you navigate:

Symbol	Meaning
*	Most Important - High priority for exam
8	Key Concept - Important insight or principle
<b>©</b>	Exam Tip - Specific advice for the exam
<u> </u>	Common Mistake - Watch out for this error
ii	Example/Calculation - Worked example
0	External Resource - Link to additional material
	Checklist Item - Something to verify you know
<b>&gt;</b>	Practice Problem - Try this yourself
4	•

# **L** PART 1: QUICK REVIEW

# <a name="review"></a>1. Quick Review: Lecture 2 Concepts

Purpose: Refresh your memory on prerequisite concepts from Lecture 2 that are essential for understanding evaluation metrics.

# **Key Models You Should Know:**

# **RSJ Model (Robertson & Spark Jones)**

- Ranks by probability of relevance: p(rel=1|q,d)
- Uses Bayes rule but doesn't leverage TF information
- Relies on relevance judgments

#### **BM25 Model**

- Approximates 2-Poisson model
- Formula:

# $c_i^BM25(tf_i) \approx log(N/df_i) \times [tf_i(k_1 + 1)] / [k_1(1 - b + b|d|/avgdl) + tf_i]$

Parameters:  $b = 0.75, k \ 1 \in [1.2, 2.0]$ 

# Language Model-Based Retrieval

Ranks by probability of generating query from document: p(q|d)

• Dirichlet smoothing formula provided in slides



# PART 2: EVALUATION METRICS (MOST IMPORTANT)

# <a name="metrics"></a>2. Evaluation Metrics

▲ Critical for Exam: This section contains the most heavily tested material. Practice all calculations by hand!

#### 2.1 Precision and Recall

#### **Definitions:**

- **Precision** = (# relevant AND retrieved) / (# retrieved)
- **Recall** = (# relevant AND retrieved) / (# relevant)

Example from Lecture: Query: "cs 589 stevens"

- Top 3 results: 2 relevant, 1 not relevant
- Precision@3 = 2/3

**Key Point:** Precision measures accuracy of what you returned; Recall measures completeness of retrieval.

# 2.2 Mean Average Precision (MAP)

#### **Step-by-Step Calculation:**

- 1. Find positions of relevant documents in your ranking (K<sub>1</sub>, K<sub>2</sub>, ... K R)
- 2. Calculate Precision@K for each relevant position
- 3. Average these precision values
- 4. Divide by total # of relevant documents (NOT # retrieved)

#### Formula:

```
AveP = \Sigma(P(k) \times rel(k)) / (number of relevant documents)
```

where the sum is over all retrieved documents n

### **Worked Example from Slides:**

Ranking #1: + + - - + - - + - -

• Relevant docs at positions: 1, 2, 5, 8

- P(a)1 = 1/1 = 1.0
- P(a)2 = 2/2 = 1.0
- P(a)5 = 3/5 = 0.6
- P(a)8 = 4/8 = 0.5
- Average Precision = (1.0 + 1.0 + 0.6 + 0.5) / 5 = 0.62
  - Note: Divided by 5 because there are 5 total relevant documents

Ranking #2: - + - - + - - -

- Relevant docs at positions: 2, 5, 7
- P(a)2 = 1/2 = 0.5
- P(a)5 = 2/5 = 0.4
- P(a)7 = 3/7 = 0.43
- Average Precision = (0.5 + 0.4 + 0.43) / 5 = 0.266

Mean Average Precision = (0.62 + 0.266) / 2 = 0.443

**©** Common Mistake: Don't divide by number of retrieved documents! Always divide by total number of relevant documents.

# 2.3 Mean Reciprocal Rank (MRR)

Use Case: When users only need ONE relevant document (e.g., navigational queries)

#### Formula:

RR = 1 / (rank of first relevant document) MRR = average of RR across all queries

### **Example from Slides:**

- Ranking #1: First relevant at position  $1 \rightarrow RR = 1/1 = 1.0$
- Ranking #2: First relevant at position  $2 \rightarrow RR = 1/2 = 0.5$
- MRR = (1.0 + 0.5) / 2 = 0.75

▲ Note: In the slides, they use  $(RR = 1.0 / (1.0 + rank_1))$  where rank starts from 0, so adjust based on indexing convention.

# 2.4 Normalized Discounted Cumulative Gain (NDCG)

Use Case: When you have multiple levels of relevance (not just binary)

• Example: 2 = click > 10s, 1 = click < 10s, 0 = no click

### **Step 1: Calculate DCG**

Two formulas (use the one your professor prefers):

#### Formula 1:

```
DCG_p = \Sigma(rel_i / log<sub>2</sub>(i+1)) for i=1 to p
```

#### Formula 2:

```
DCG_p = \Sigma((2^{i-1}) / \log_2(i+1)) for i=1 to p
```

# **Step 2: Calculate IDCG (Ideal DCG)**

- Sort relevance scores in descending order
- Calculate DCG for this ideal ranking

# **Step 3: Calculate NDCG**

```
NDCG = DCG / IDCG
```

# **Worked Example from Slides:**

Ranking #1: [2, 0, 1, 2, 2, 1, 0, 0, 0, 2]

# DCG@4 using Formula 1:

```
= 2/\log_2(2) + 0/\log_2(3) + 1/\log_2(4) + 2/\log_2(5)
= 2/1 + 0 + 1/2 + 2/2.32
= 2 + 0 + 0.5 + 0.86
= 3.3613
```

# **IDCG**(a)**4:** Ideal ranking = [2, 2, 2, 2]

```
= 2/\log_2(2) + 2/\log_2(3) + 2/\log_2(4) + 2/\log_2(5)
= 2 + 1.26 + 1 + 0.86
= 5.1232
```

#### NDCG@4 = 3.3613 / 5.1232 = 0.656

**Important:** If # relevant items < k, IDCG calculation stops at the number of relevant items!



# <a name="methodologies"></a>3. Evaluation Methodologies

**Focus:** Understanding how IR systems are evaluated in practice - from lab experiments to real-world A/B testing.

# 3.1 The Cranfield Experiment (1958)

## **Basic Ingredients:**

- 1. A corpus of documents (~1.4k paper abstracts)
- 2. A set of queries (225 queries)
- 3. Binary relevance judgments for each (query, document) pair
- 4. Reusable relevance judgments

Key Innovation: Created a standardized test collection for IR evaluation

# 3.2 Pooling Strategy

**Problem:** Too many (query, document) pairs to annotate

•  $225 \text{ queries} \times 1,400 \text{ docs} = 315,000 \text{ pairs!}$ 

#### **Solution:**

- 1. Run K different systems
- 2. For each system, take top 100 results
- 3. Pool (union) all these documents
- 4. Only annotate documents in this pool

Why it works: Relevant documents likely appear in top results of at least one system

**©** Exam Tip: Understand that pooling reduces annotation cost while maintaining quality

# 3.3 Text REtrieval Conference (TREC)

#### **Key Facts:**

- Since 1992, hosted by NIST
- Uses Cranfield methodology with pooling
- Relevance judgments based on human annotations
- Goes beyond keyword matching
- Different tracks: Web, Question Answering, Microblog, etc.

#### **Sample TREC Query Format:**

```
<top>
<num> Number: 794
<title> pet therapy
<desc> Description:
How are pets or animals used in therapy for humans and what are the benefits?
<narr> Narrative:
Relevant documents must include details of how pet- or animal-assisted
therapy is or has been used...

</top>
```

# 3.4 Online Evaluation & A/B Testing

# Why Online Evaluation?

• TREC-style: Explicit, difficult to scale, gets outdated

• Click logs: Implicit, large-scale, up-to-date

# **Key Concepts:**

#### **Position Bias**

• Higher positions get more attention

• Same item gets fewer clicks in lower positions

• Models:

• Baseline: No position bias

• Mixture: Relevance + constant bias

• Cascade: Linear traversal; documents below clicked result not examined (best for lower ranks)

#### **Cascade Model Formula:**

```
c_di = r_d \prod (1 - r_{docinrank:j}) for j=1 to i-1
```

### **User Click Logs**

### **System logs track:**

- Timestamp
- Session ID
- Query ID and content

- Items viewed (sequential order)
- Click/no-click for each item
- User demographics, history, location, device
- Dwell time, browsing time
- Eye tracking information

# **Click Logs Storage:**

- Stored in large tables
- Extract subsets using SQL queries

# **Decoy Effects**

User preferences change based on context:

#### **Example:**

- Option A (\$400, 20G) vs Option B (\$500, 30G): 50-50 split
- Add decoy (\$550, 20G): Now Option B gets 60% preference

Implication: Can't assume relevance is fixed; it's context-dependent

#### A/B Testing

#### **Process:**

- 1. Split traffic: 50% System A, 50% System B
- 2. Compare metrics (clicks, retention, revenue)
- 3. Use statistical significance tests

#### **Statistical Tests:**

## **Sign Test:**

- Counts how many queries System A beats System B
- Simple but ignores magnitude
- Python: (statsmodels.stats.descriptivestats.sign test)

#### **Wilcoxon Test:**

- Considers both direction AND magnitude
- Formula:  $W = \Sigma[\operatorname{sgn}(x_{2,i} x_{1,i}) \cdot R_i]$
- More powerful than sign test

• Python: (scipy.stats.wilcoxon)

# p-value Interpretation:

- p < 0.05: Statistically significant (reject null hypothesis)
- $p \ge 0.05$ : Not significant (can't conclude difference)

#### **Multi-Armed Bandit**

#### **Problem with A/B Testing:**

• Fixed 50/50 split means poor system hurts user experience

#### **Solution:**

- Dynamically adjust traffic based on performance
- Explore vs Exploit tradeoff

# **Upper Confidence Bound Formula:**

$$\hat{\mu}_i(t-1) + \sqrt{(2 \log t / T_i(t-1))}$$

#### Where:

- $\mu$  i(t-1): Estimated mean reward for arm i
- T i(t-1): Number of times arm i has been pulled
- t: Current time step

**Result:** Better system gets more traffic over time while still exploring alternatives

# **Interleaving**

# Method:

- 1. Merge results from System A and System B
- 2. Remove duplicates
- 3. Show combined ranking to user
- 4. Track which system's documents get clicked

# **Example:**

```
System A: [Kernel machines, Kernel machines, SVMs, SVM-light, ...]

System B: [Kernel machines, SVMs, Intro to SVMs, Lucent SVM demo, ...]

Interleaved (duplicates removed):
[Kernel machines (A), Kernel machines (B), SVMs (A), SVM-light (A),
Intro to SVMs (B), Lucent SVM demo (B), ...]

If clicks: Kernel machines (A), SVMs (A), SVM-light (B), Intro to SVMs (B)

→ A clicks = 3, B clicks = 1
```

### Advantages:

- More sensitive than A/B testing
- Same user sees both systems
- Controls for user variability

# 🔁 PART 4: RELEVANCE FEEDBACK

# <a name="feedback"></a>4. Relevance Feedback

**Key Idea:** Using user interactions to improve search results through query refinement.

#### 4.1 Motivation

**Problem:** Users don't always know how to express their information need

Example: Query "best phone"

- Does user prefer lower-priced or high-end?
- Larger storage or better camera?

Solution: Learn from user interactions (clicks, dwell time) to refine the query

# 4.2 Rocchio Feedback (Vector Space Model)

# Formula:

```
q_F = \alpha q + (\beta/|D_r|)\Sigma d_r - (\gamma/|D_n|)\Sigma d_n
```

#### Where:

q: Original query vector

- q F: Feedback-adjusted query
- D r: Set of relevant documents
- D n: Set of non-relevant documents
- α: Weight for original query
- $\beta$ : Weight for relevant docs (typically  $\beta >> \gamma$ )
- γ: Weight for non-relevant docs

#### **Intuition:**

- Moves query vector toward relevant documents
- Moves query vector away from non-relevant documents
- Positive evidence weighted more than negative

# Visual Representation:

Before feedback: q (original query)

After feedback: q\_F (moved toward relevant docs cluster)

#### **Practical Issues:**

- 1. Large vocabularies Only consider important words
- 2. Requires explicit feedback User must label docs
- 3. Robust and effective when feedback available
- Interactive Demo: <a href="https://tinyurl.com/4bkw7cj2">https://tinyurl.com/4bkw7cj2</a>

# 4.3 Pseudo-Relevance Feedback

**Problem:** What if we don't have relevance judgments?

Solution: Assume top-k retrieved documents are relevant

#### **Process:**

- 1. Run initial query
- 2. Retrieve top k documents (e.g., k=10)
- 3. Treat these as "pseudo-relevant"
- 4. Apply feedback algorithm
- 5. Re-rank with new query

# Why It Works:

Query: "fish tank"

Top results contain: "aquarium", "filter", "water", "gravel", "fish"

Expand query: "fish tank aquarium filter water"

Result: Better recall, finds documents using different terminology

# For Vector Space Model:

- Use Rocchio with top-k docs as D\_r
- D\_n can be empty or lower-ranked docs

# For Language Model:

```
\theta_{q} = \lambda \theta_{q} + (1 - \lambda)\theta_{d}
```

Where  $\theta_d$  is estimated from feedback documents

# 4.4 Feedback Language Model

#### Formula:

$$\theta \land FB = \lambda \theta_q + (1-\lambda)\theta_d$$

## **Components:**

- θ\_q: Query language model
- $\theta$  d: Document language model (from relevant docs)
- λ: Interpolation parameter (controls mixing)

# **Process:**

- 1. Get initial query model  $\theta$  q
- 2. Retrieve documents
- 3. Estimate  $\theta$  d from top documents
- 4. Combine using mixture model
- 5. Re-rank using  $\theta$ ^FB

# Estimating $\theta$ d:

- Can use EM algorithm
- Accounts for query-specific vs general terms

# **Quiz Question from Lecture 3:**

Given feedback documents, calculate  $\theta$ ^FB:

 $\theta ^{FB} = \lambda \theta_{q} + (1-\lambda)\theta_{d}$ 

If  $\lambda = 0.7$ , this means:

- 70% weight to original query
- 30% weight to feedback documents

**Result (from quiz):** Probability of "airport security" = **0.38** 

# 4.5 Query Expansion/Reformulation

# **Techniques:**

#### 1. Manual Thesaurus:

- Use WordNet or domain-specific thesaurus
- Example: "skin itch" → add "pruritus", "integumentary system"

#### 2. Automatic Thesaurus:

- Learn word similarities from corpus
- Example nearest neighbors:
  - "absolutely"  $\rightarrow$  "absurd", "whatsoever", "totally"
  - "captivating" → "shimmer", "stunningly", "superbly"

# 3. Query Log Mining:

- Analyze what users searched and clicked
- Example: Users who search "movie tickets" also search "showtimes"

#### **Google's Query Expansion:**

- "what is the most..." → autocomplete suggestions
- "yoga mat" → filter by price, brand



# **PART 5: PRACTICE PROBLEMS**

# <a name="practice"></a>5. Practice Problems

**Instructions:** Try solving these problems by hand before revealing the solutions. These mirror the quiz and exam format.

# **Problem 1: Calculate MAP**

Given two systems with rankings for 2 queries (5 relevant docs total each):

# Query 1:

- System A: (++-+-+-+) (+ = relevant)
- System B: (+-+-+-+)

# Calculate MAP for each system.

<details> <summary>Click to see solution</summary>

# System A (Query 1):

- Relevant at positions: 1, 2, 4, 6, 10
- P@1 = 1/1 = 1.0
- P(a)2 = 2/2 = 1.0
- P(a)4 = 3/4 = 0.75
- P(a)6 = 4/6 = 0.67
- P(a)10 = 5/10 = 0.5
- AP = (1.0 + 1.0 + 0.75 + 0.67 + 0.5) / 5 = 0.784

# System B (Query 1):

- Relevant at positions: 1, 3, 5, 8, 10
- P(a)1 = 1/1 = 1.0
- P(a)3 = 2/3 = 0.67
- P(a)5 = 3/5 = 0.6
- P@8 = 4/8 = 0.5
- P(a)10 = 5/10 = 0.5
- AP = (1.0 + 0.67 + 0.6 + 0.5 + 0.5) / 5 = 0.654

For full MAP, average across both queries.

</details>

# Problem 2: Calculate NDCG@5

**Ranking:** [3, 2, 0, 1, 2]

Relevance scale: 0-3

<details> <summary>Click to see solution</summary>

# **Using Formula 1:**

# DCG@5:

```
= 3/\log_2(2) + 2/\log_2(3) + 0/\log_2(4) + 1/\log_2(5) + 2/\log_2(6)
= 3/1 + 2/1.585 + 0/2 + 1/2.322 + 2/2.585
= 3 + 1.262 + 0 + 0.431 + 0.774
= 5.467
```

# **IDCG@5:** Ideal = [3, 2, 2, 1, 0]

```
= 3/1 + 2/1.585 + 2/2 + 1/2.322 + 0
= 3 + 1.262 + 1 + 0.431 + 0
= 5.693
```

# NDCG@5 = 5.467 / 5.693 = 0.96

</details>

# **Problem 3: Calculate MRR**

# 3 queries:

- Query 1: First relevant at position 1
- Query 2: First relevant at position 3
- Query 3: First relevant at position 2

<details> <summary>Click to see solutionSummary> ``` RR1 = 1/1 = 1.0 RR2 = 1/3 = 0.333 RR3 = 1/2 = 0.5

MRR = (1.0 + 0.333 + 0.5) / 3 = 0.611

```
---
### Problem 4: Feedback Language Model (From Quiz)

**Scenario:**
- Query: "airport security"
- λ = 0.7
- θ_q: probability of "airport" in query = 0.5
- θ_d: probability of "airport" in feedback docs = 0.2

**Calculate θ^FB for "airport":**

<details>
<summary>Click to see solution</summary>
```

```
\theta FB = \lambda \theta_q + (1-\lambda)\theta_d
= 0.7 × 0.5 + 0.3 × 0.2
= 0.35 + 0.06
= 0.41
```

```
**Answer from quiz:** 0.38 (may use different values or formula variant)

</details>

**Question:** If you have only 3 relevant items but calculate NDCG@5, how do you compute IDCG?

<details>
<summary>Click to see solution</summary>

**Scenario:** Relevant items have scores [2, 1, 1]

**IDCG@5 calculation:**
- Ideal ranking: [2, 1, 1, 0, 0]
- Only compute up to position 3 (last relevant item):
```

IDCG@5 = 
$$2/\log_2(2) + 1/\log_2(3) + 1/\log_2(4)$$
  
=  $2/1 + 1/1.585 + 1/2$ 

$$= 2 + 0.631 + 0.5$$

= 3.131

```
**Note: ** Positions 4 and 5 contribute 0, so can be ignored.
***Impact:*** nDCG value will be lower if you don't retrieve all relevant items in top-k.
</details>
#  PART 6: ADDITIONAL RESOURCES
## <a name="resources"></a>6. Additional Resources
> **Enhance Your Learning: ** Curated videos, readings, and tools to deepen your understanding.
### **Recommended Reading**
1. **Stanford IR Book** (Your professor's reference)
 - Chapter 8: Evaluation in Information Retrieval
 - https://nlp.stanford.edu/IR-book/
 - Sections 8.1-8.4 (TREC, Precision/Recall)
 - Section 8.7 (Results Snippets)
2. **Manning, Raghavan, Schütze - Introduction to Information Retrieval**
 - Chapter 9: Relevance Feedback
 - Free online: https://nlp.stanford.edu/IR-book/pdf/09expand.pdf
3. **Research Papers:**
 - Järvelin & Kekäläinen (2002): "Cumulated gain-based evaluation of IR techniques" (NDCG)
 - Craswell et al. (2008): "An experimental comparison of click position-bias models"
### ## **Video Tutorials**
#### Evaluation Metrics
1. **Precision & Recall (StatQuest):**
 - https://www.youtube.com/watch?v=vP06aMoz4v8
 - Duration: 15 minutes
 - Great visual explanations
```

```
2. **MAP Explained:**
 - https://www.youtube.com/watch?v=vYVSVhYNX7E
 - Duration: 10 minutes
 - Step-by-step calculation
3. **NDCG Intuition:**
 - https://www.youtube.com/watch?v=7L76BaqOYCI
 - Duration: 12 minutes
 - Why DCG uses logarithmic discount
4. **F1 Score (related to Precision/Recall):**
 - https://www.youtube.com/watch?v=jJ7ff7Gcq34
 - Duration: 8 minutes
#### Online Evaluation
5. **A/B Testing Basics:**
 - https://www.youtube.com/watch?v=zFMgpxG-chM
 - Duration: 20 minutes
 - Statistical testing fundamentals
6. **Multi-Armed Bandit:**
 - https://www.youtube.com/watch?v=e3L4VocZnnQ
 - Duration: 18 minutes
 - Exploration vs exploitation
#### Mathematical Background
7. **Logarithms Review (Khan Academy):**
 - https://www.khanacademy.org/math/algebra2/x2ec2f6f830c9fb89:logs
 - Useful for understanding DCG discount factor
8. **Statistical Significance:**
 - https://www.youtube.com/watch?v=5Z9OIYA8He8
 - Duration: 15 minutes
 - p-values and hypothesis testing
### ** **Practice Tools & Code**
#### Python Libraries
```python
# Precision, Recall, F1
from sklearn.metrics import precision_score, recall_score, fl_score
# Statistical tests
```

```
from scipy.stats import wilcoxon
from statsmodels.stats.descriptivestats import sign_test

# Calculate NDCG
from sklearn.metrics import ndcg_score
```

# **Example Code: Calculate MAP**

```
python
def average_precision(relevant_positions, total_docs):
  Calculate average precision for a single query.
  Args:
     relevant_positions: List of positions (1-indexed) where relevant docs appear
     total_docs: Total number of relevant documents
  Returns:
     Average precision score
  if not relevant positions or total docs == 0:
     return 0.0
  precisions = []
  for i, pos in enumerate(sorted(relevant_positions)):
     precision_at_k = (i + 1) / pos
     precisions.append(precision_at_k)
  return sum(precisions) / total_docs
# Example
relevant_pos = [1, 2, 5, 8] # relevant docs at positions 1, 2, 5, 8
total relevant = 5 # 5 relevant docs total
ap = average_precision(relevant_pos, total_relevant)
print(f'Average Precision: {ap:.3f}")
```

# **Example Code: Calculate NDCG**

python

```
import numpy as np
def dcg at k(relevances, k):
  """Calculate DCG@k"""
  relevances = np.asarray(relevances)[:k]
  if relevances.size:
     \#Formula: rel_i / log2(i+2) for i starting at 0
     return np.sum(relevances / np.log2(np.arange(2, relevances.size + 2)))
  return 0.0
def ndcg_at_k(relevances, k):
  """Calculate NDCG@k"""
  dcg = dcg_at_k(relevances, k)
  idcg = dcg_at_k(sorted(relevances, reverse=True), k)
  if ideg == 0:
     return 0.0
  return dcg / idcg
# Example
ranking = [3, 2, 0, 1, 2]
k = 5
ndcg = ndcg_at_k(ranking, k)
print(f"NDCG@{k}: {ndcg:.3f}")
```

#### **Interactive Tools**

#### 1. Rocchio Feedback Demo:

- <a href="https://tinyurl.com/4bkw7cj2">https://tinyurl.com/4bkw7cj2</a>
- Visualize query refinement

#### 2. IR Metrics Calculator:

- Create your own spreadsheet to practice
- Or use online calculators (search "MAP calculator IR")

# 3. Python Notebook Examples:

- Google Colab notebook for IR metrics
- Practice with your own data

# **II** Comparison Tables

When to Use Which Metric?

Metric	Use When	Pros	Cons
Precision@k	Fixed cutoff important	Simple, interpretable	Ignores rank within top-k
Recall@k	Need to find all relevant	Measures completeness	Doesn't penalize bad ranking
MAP	Multiple relevant docs, order matters	Emphasizes top ranks	Binary relevance only
MRR	Only need 1 relevant doc	Good for navigational queries	Ignores other relevant docs
NDCG	Graded relevance, order matters	Handles multiple relevance levels	More complex to calculate
4	'	•	▶

# **Evaluation Methodology Comparison**

Method	Pros	Cons	Best For
TREC-style (Pooling)	Reusable, controlled	Expensive, gets outdated	Research, benchmarking
Click logs	Large-scale, up-to-date	Biased, noisy	Industry, real-world
A/B testing	Real user feedback	Needs traffic, statistical power	Production systems
Interleaving	More sensitive than A/B	More complex	Fine-tuning systems
<b>▲</b>	•	•	▶

# Quick Reference Sheet

**Print This Page:** This section is designed to be printed and used as a quick reference during your final review.

# **Formulas Summary**

```
Precision@k = (\# relevant in top k) / k
Recall@k = (\# relevant in top k) / (total \# relevant)
F1@k = 2 \times (Precision@k \times Recall@k) / (Precision@k + Recall@k)
Average Precision = \Sigma(P(k) \times rel(k)) / (\# relevant docs)
              where k ranges over all positions
MAP = (1/|Q|) \Sigma AveP(q) for all queries Q
MRR = (1/|Q|) \Sigma (1 / rank of first relevant doc)
DCG_p = \Sigma (rel_i / log<sub>2</sub>(i+1)) [i from 1 to p]
Alternative DCG_p = \Sigma ((2^rel_i - 1) / log<sub>2</sub>(i+1))
NDCG_p = DCG_p / IDCG_p
Rocchio: q_F = \alpha q + (\beta/|D_r|)\Sigma d_r - (\gamma/|D_n|)\Sigma d_n
Feedback LM: \theta^FB = \lambda\theta_q + (1-\lambda)\theta_d
Upper Confidence Bound: \mu_i(t-1) + \sqrt{2 \log t / T_i(t-1)}
```

# **Common Parameter Values**

Rocchio:  $\alpha = 1.0$ ,  $\beta = 0.75$ ,  $\gamma = 0.15$ BM25: b = 0.75,  $k_1 \in [1.2, 2.0]$ Feedback LM:  $\lambda \in [0.5, 0.9]$ Pseudo-relevance:  $k \in [5, 20]$  documents

# PRE-EXAM CHECKLIST & FINAL REVIEW

# **✓** Pre-Exam Checklist

Use This: Go through this checklist 24 hours before your exam to identify any gaps in your knowledge.

# **Calculation Skills**

- Can calculate Precision@k and Recall@k by hand
- Can calculate Average Precision for a single query (remember: divide by total # relevant!)

Can calculate MAF across multiple queries
Can calculate MRR for navigational queries
Can calculate DCG@k using the logarithmic discount formula
Can determine IDCG@k (remember: sort by relevance descending!)
Can calculate NDCG@k from DCG and IDCG
Can apply Rocchio feedback formula
Conceptual Understanding
Understand when to use each evaluation metric
■ Know the difference between Precision and Recall
☐ Can explain why MAP divides by # relevant, not # retrieved
Understand why NDCG uses logarithmic discounting
Can explain the Cranfield experiment methodology
Understand how pooling reduces annotation cost
■ Know the three types of position bias models (Baseline, Mixture, Cascade)
Can explain the difference between A/B testing and interleaving
Understand how Rocchio feedback works (move toward relevant, away from non-relevant)
Can explain pseudo-relevance feedback and its risks
Common Pitfalls to Avoid
Common Pitfalls to Avoid  MAP: Don't divide by # retrieved documents - always use total # relevant
■ MAP: Don't divide by # retrieved documents - always use total # relevant
■ MAP: Don't divide by # retrieved documents - always use total # relevant ■ NDCG: Don't forget the logarithmic discount - it's log₂(i+1), not just (i+1)
■ MAP: Don't divide by # retrieved documents - always use total # relevant ■ NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1) ■ IDCG: When # relevant < k, IDCG only sums up to # relevant items
■ MAP: Don't divide by # retrieved documents - always use total # relevant ■ NDCG: Don't forget the logarithmic discount - it's log₂(i+1), not just (i+1) ■ IDCG: When # relevant < k, IDCG only sums up to # relevant items ■ MRR: Only considers the FIRST relevant document
<ul> <li>MAP: Don't divide by # retrieved documents - always use total # relevant</li> <li>NDCG: Don't forget the logarithmic discount - it's log₂(i+1), not just (i+1)</li> <li>IDCG: When # relevant &lt; k, IDCG only sums up to # relevant items</li> <li>MRR: Only considers the FIRST relevant document</li> <li>Indexing: Check if positions start at 0 or 1 (be consistent!)</li> </ul>
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)  Pseudo-relevance: Can make things worse if initial results are bad
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)  Pseudo-relevance: Can make things worse if initial results are bad  Exam Strategy
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log2(i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)  Pseudo-relevance: Can make things worse if initial results are bad  Exam Strategy  1. Read problems carefully - Check if positions are 0-indexed or 1-indexed
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)  Pseudo-relevance: Can make things worse if initial results are bad  Exam Strategy  1. Read problems carefully - Check if positions are 0-indexed or 1-indexed  2. Show your work - Partial credit for correct process
MAP: Don't divide by # retrieved documents - always use total # relevant  NDCG: Don't forget the logarithmic discount - it's log <sub>2</sub> (i+1), not just (i+1)  IDCG: When # relevant < k, IDCG only sums up to # relevant items  MRR: Only considers the FIRST relevant document  Indexing: Check if positions start at 0 or 1 (be consistent!)  Rocchio: β should be >> γ (positive evidence weighted more)  Pseudo-relevance: Can make things worse if initial results are bad  Exam Strategy  1. Read problems carefully - Check if positions are 0-indexed or 1-indexed  2. Show your work - Partial credit for correct process  3. Double-check denominators - MAP uses # relevant, not # retrieved

# **©** Final Tips for Success

# **Study Strategy**

1. Practice calculations by hand first

- Don't rely only on code/calculators
- Work through examples step-by-step
- The quiz solutions are excellent practice

# 2. Understand the "why"

- Why does NDCG use log discount? (Users scan top results more carefully)
- Why  $\beta > \gamma$  in Rocchio? (Positive evidence more reliable)
- Why does pooling work? (Good systems will find relevant docs)

# 3. Review quiz solutions thoroughly

- Quiz closely reflects exam format
- Pay attention to edge cases
- Note the exact formulas used

#### 4. Create your own examples

- Make up small ranking examples
- Calculate metrics by hand
- Verify with Python code

# **Time Management**

- Calculation problems: Budget ~10 min each
- Conceptual questions: ~5 min each
- Leave 10-15 min for review

# **During the Exam**

### **Do:**

- Write down formulas first
- Show all calculation steps
- Box or highlight final answers
- Check units and ranges (e.g., NDCG  $\in$  [0,1])

#### X Don't:

- Rush through calculations
- Forget to label your answers
- Skip showing work (even if answer is wrong, process matters!)

# **External Links Summary**

#### **Recommended Videos:**

- Precision/Recall: <a href="https://www.youtube.com/watch?v=vP06aMoz4v8">https://www.youtube.com/watch?v=vP06aMoz4v8</a>
- MAP: <a href="https://www.youtube.com/watch?v=vYVSVhYNX7E">https://www.youtube.com/watch?v=vYVSVhYNX7E</a>
- NDCG: <a href="https://www.youtube.com/watch?v=7L76BaqOYCI">https://www.youtube.com/watch?v=7L76BaqOYCI</a>
- A/B Testing: <a href="https://www.youtube.com/watch?v=zFMgpxG-chM">https://www.youtube.com/watch?v=zFMgpxG-chM</a>
- Multi-Armed Bandit: <a href="https://www.youtube.com/watch?v=e3L4VocZnnQ">https://www.youtube.com/watch?v=e3L4VocZnnQ</a>

### Reading:

- Stanford IR Book: <a href="https://nlp.stanford.edu/IR-book/">https://nlp.stanford.edu/IR-book/</a>
- Interactive Rocchio Demo: <a href="https://tinyurl.com/4bkw7cj2">https://tinyurl.com/4bkw7cj2</a>

#### **Tools:**

- Python sklearn: (from sklearn.metrics import ndcg score, precision score)
- SciPy: (from scipy.stats import wilcoxon)

# 🔯 Last-Minute Review (30 Minutes Before Exam)

**Emergency Prep:** If you only have 30 minutes, focus on this section!

# **Quick Formulas Review (5 min)**

```
P@k = relevant in top k / k

R@k = relevant in top k / total relevant

AP = \Sigma(P(k) \times rel(k)) / total\_relevant
DCG = \Sigma(rel\_i / log_2(i+1))
NDCG = DCG / IDCG
```

# **Practice One Problem Each (15 min)**

- 1. Calculate MAP for one ranking
- 2. Calculate NDCG@5 for one ranking
- 3. Apply Rocchio formula once

# **Concept Review (10 min)**

- Pooling: Take top-k from multiple systems, annotate only pool
- Pseudo-relevance: Assume top-k are relevant, expand query

- Position bias: Higher positions get more attention
- Interleaving: Mix two systems' results, compare clicks

# **Lesson** You've Got This!

#### What You've Mastered:

- All major evaluation metrics (P/R, MAP, MRR, NDCG)
- **V** Evaluation methodologies (TREC, pooling, online)
- Statistical testing (Sign, Wilcoxon)
- Relevance feedback (Rocchio, pseudo-relevance, LM)
- Practical considerations (position bias, A/B testing, bandits)

#### Remember:

- **!** The quiz solutions match the exam format closely
- 🐔 If you can solve those problems by hand, you're ready!
- 6 Focus on understanding WHY, not just memorizing formulas
- Manage your time during the exam
- Show your work for partial credit

#### **Final Confidence Boosters:**

- 1. You have comprehensive notes covering all testable material
- 2. You have worked examples for every type of problem
- 3. You have a formula reference sheet ready
- 4. You know the common mistakes to avoid
- 5. You have a clear exam strategy

# Good luck on your midterm! 🍀 💵

You've prepared well. Trust your preparation and stay calm during the exam!

# **Document Information**

#### **Document Details:**

Title: CS 589 Lecture 3 - Complete Study Guide

Course: CS 589 - Text Mining and Information Retrieval

**Institution:** Stevens Institute of Technology

#### **Topics Covered:**

- Evaluation Metrics (Precision, Recall, MAP, MRR, NDCG)
- Evaluation Methodologies (Cranfield, TREC, Pooling)
- Online Evaluation (A/B Testing, Interleaving, Multi-Armed Bandits)
- Relevance Feedback (Rocchio, Pseudo-relevance, Language Models)

# **Exam Preparation Materials:**

- Worked examples from lecture slides
- **Quiz** solutions with detailed explanations
- Practice problems with step-by-step solutions
- V Formula quick reference sheet
- Pre-exam checklist

#### **Additional Resources:**

- Video tutorials with timestamps
- Python code examples
- Interactive demos
- External reading materials

# © Study Guide prepared for CS 589 students

For educational purposes only

## Questions or Found an Error?

If you spot any errors or have questions about the material, please:

- 1. Review the lecture slides at the referenced page numbers
- 2. Check the quiz solutions document
- 3. Consult the Stanford IR textbook (Chapter 8)
- 4. Ask your professor or TA during office hours

